

A Bird's Eye View on Requirements Engineering and Machine learning

1st Tahira Iqbal
fortiss GmbH
Munich, Germany
iqbal@fortiss.org

2nd Parisa Elahidoost
fortiss GmbH
Munich, Germany
elahidoost@fortiss.org

3rd Levi Lúcio
fortiss GmbH
Munich, Germany
lucio@fortiss.org

Abstract—Machine learning (ML) has demonstrated practical impact in a variety of application domains. Software engineering is a fertile domain where ML is helping in automating different tasks. In this paper, our focus is the intersection of software requirement engineering (RE) and ML. To obtain an overview of how ML is helping RE and the research trends in this area, we have surveyed a large number of research articles. We found that the impact of ML can be observed in requirement elicitation, analysis and specification, validation and management. Furthermore, in these categories, we discuss the specific problem solved by ML, the features and ML algorithms used as well as datasets, when available. We outline lessons learned and envision possible future directions for the domain.

Index Terms—Requirements Engineering, Machine learning, State of the art, Overview

I. INTRODUCTION

¹ Machine learning algorithms have been shown to have considerable practical importance in many application domains. This is especially true of domains where large databases are available and a need for exploring some kind of consistency exists or, domains where a program needs to adapt itself to changes [55]. Requirements engineering is a critical part of software engineering and it seems appropriate to use machine learning methods for requirements engineering tasks. Because requirements specification documents are mainly given in natural language, ML can be useful by emulating human processing.

This paper aims to present a survey of how ML benefits existing RE approaches. More precisely, we pursue the following research questions:

RQ1: What is the current state of the practice in ML & RE?

RQ2: What types of learning methods are used when ML is applied to RE?

RQ3: Which are the RE problems that currently use ML methods?

RQ4: Is using ML methods improving RE?

To reply to these research questions we have performed a literature review, split into data preparation, data collection, and data analysis phase. First, a search string was prepared based on the research questions, then a search was performed over a predefined set of databases and all identified studies were assessed by means of title and abstract. Our literature

review is not meant to be an exhaustive study of the field – rather we are offering a snapshot of the current state-of-the-art by borrowing some techniques from Systematic Literature Reviewing.

The major contributions of this article are as follows:

- We provide an overview of the current state of the art of some of the challenges RE faces that may be handled through ML techniques. We focus on two important aspects:
 - Providing an overview of the ML problem categories (classification, regression, clustering, etc.) in use for the support of RE tasks (elicitation, analysis, validation and management).
 - Providing an overview of the common ML models (decision tree, K-Nearest Neighbors, Naive Bayesian, etc.) for tackling RE problems and the data sets if available.
- We analyze the literature to discover trends and lessons on the use of ML in RE.

The paper is organized as follows. The rest of this section provides background information on ML and RE required to understand the remaining of the paper. Section II provides an overview of RE tasks where ML has been used. In section III we summarize the major findings of our study. Finally, in section IV we state the threats to the validity of our work and section V concludes the paper.

A. Machine Learning

Machine-Learning (ML) [33] is a range of algorithms to approximate discover patterns in data. Historically, models and heuristics are human-built exhaustive prescriptions of how a system should behave. ML is grounded on different premises: rather than relying on humans to input all the possible cases the system can handle, the field attempts to extrapolate patterns from a representative set of examples that illustrates expected behaviors. The way in which a learning algorithm operates attempts to emulate the way in which humans learn: from a set of examples, a general model for a behavior is induced.

Many learning algorithms exist, based on different visions of how learning happens in practice [13]. All these algorithms have in common the notion of *features*. Features correspond to characteristics of what is being learned and provide the

¹This work was supported in part by the European Union's Horizon 2020 Research and Innovation program under grant agreement no. 674875

grounds for the algorithm to abstract from the complexities of the real world. Assume for example that an algorithm should learn, based on a brain scan of a medical patient, to decide whether that patient has brain cancer or not. A number of *features* such as for example the “number of irregular objects in the scan”, the “color of such objects”, the “disposition of such objects” would be provided to the algorithm. Additionally, the algorithm is provided with a number of brain scans together with annotations that summarize decisions previously taken on them (in our example cancer found / cancer not found). Such datasets are called the *training data*. The learning algorithm then undergoes a *training phase*. It attempts to find an internal model that allows it to map the decisions to the brain scans, given the training data. The model obtained from the training step is useful if it performs well (generalizes) when applied to new data from outside the training set – in our example, when it can accurately diagnose brain cancer for new brain scans. Such generalization is based on the premise that inputs that are “closer”, in terms of the given *features*, should lead to “closer” outputs.

The established literature in the domain (e.g. [33]) typically considers three types of machine learning:

1) *Supervised learning*: consists of learning a function using training data including annotations of the outcome of the function to be learned (e.g. patient John Doe with a certain number of physiological characteristics was diagnosed with cancer). Supervised learning can be roughly subdivided in two popular problems: *classification* and *regression*. When the output of the function being learned is composed of categorical values (i.e. classes), then we have a classification problem. The goal is to learn how to link instances or samples for a number of parameters to a certain class of values (e.g. healthy patient or unhealthy patient). However, if the co-domain of the function being learned contains continuous values, then we face a regression problem (e.g. predict the body temperature of a patient given some clinical features of the patient).:

2) *Unsupervised learning*: In some cases, the output of the function being learned is not given and we have to find patterns in the training data “blindly.” This is called an *unsupervised learning* problem. For instance, one may want to cluster patients based on symptoms they have.:

3) *Reinforcement Learning*: can be seen as an intermediate problem the co-domain of the function being learned is not given but the procedure is guided nevertheless. In reinforcement learning, an agent has to find a sequence of actions leading to a success. The fact that the sequence leads to a success, is not known in advance, but rewards are given to the agent in order for it to know if it follows a path to success. Typical examples can be found in gaming, where an agent receives a reward when it wins the game. From the different chains of actions that led him to a reward, the agent must generalize to find how to improve at that same game.:

B. Requirements Engineering

Software systems are developed over millions of lines of code, number of modules and documents. The primary goal

of the software system is to satisfy users by developing the software that can meet their needs and expectations. This goal is achievable by applying different methodologies and engineering techniques. One of the key factor is to understand and identify the needs of users, also known as, software requirements. Software requirement engineering is the process that helps to determine the requirements in a systematic way to know what functionalities the targeted system should have to fulfil user’s needs. Formally RE is defined as [54]:

“Requirements engineering is the branch of software engineering concerned with the real-world goals for, functions of, and constraints on software systems. It is also concerned with the relationship of these factors to precise specifications of software behavior, and to their evolution over time and across software families.”

Software requirements plays a key role in the success of a project. In the USA, a survey was conducted over 8380 projects by 350 companies to know the project failure rates. The report [45] results showed only 16.2% projects were completed successfully and one-half (52.7%) of projects met with challenges and were completed with partial functionalities, time delays and over budget. Almost 31% of the projects were never completed. The main cause identified by executive managers was poor requirements engineering. The main culprits were the lack of user involvement (13%), requirements incompleteness (12%), changing requirements (11%), unrealistic expectations (6%), and unclear objectives (5%).

Software requirement engineering has mainly four phases; *requirement elicitation*, *requirement analysis*, *requirement documentation* and *requirement verification* [27]. Requirement elicitation [8], [56] helps to understand the stakeholders needs, e.g. what features he wants in the software. Requirement elicitation techniques are mostly derived by the social sciences, organizational theory, knowledge engineering and practical experience. For requirements elicitation, different techniques exist in the literature that include interviews, questionnaires and ethnography etc. to understand the stakeholder requirements. Requirement analysis [35] is the next step after requirement elicitation. In this phase, software requirements are analyzed to check for conflicts and consistency of requirements. It also makes sure that the requirements are clear and complete. Furthermore, the agreed requirements are documented. This documentation has a clear and precise definition of the system functionalities. It also acts as an agreement between stakeholders and developers. These requirements are documented usually as diagrams, mathematical formulae or natural languages. These documents are used and iterated upon until the end of the projects.

System requirements are classified into business requirements, user requirements, functional requirements (FR) and non-functional requirements (NFR). Functional requirements are the system requirements that include the main features and characteristics of the desired system. Non-functional require-

ments are the system properties and constraints [9], [19]. NFRs set the criteria for judging the operation of the system e.g. performance, availability or reliability. Business requirements are specified to address business objectives, vision, and goals. It is defined at a high level to keep the knowledge from organization or company side for designing the products. User requirements are the wish list for the system from users. User requirements are valuable for ensuring that system performs similarly as users wanted it to do.

C. Text Prepration for ML

Requirements are written in the natural language such as documents or reviewing some application. These can be in a variety of forms from a list of individual words, sentences, multiple paragraphs, short texts with special characters or others. Before applying a machine learning algorithm on them different steps employed to transform words into features such as text mining, and natural langugae processing (NLP). NLP text preprocessing phase relies majorly on pre-built dictionaries, databases, and rules. The common preprocessing steps in our literature include capitalization, tokenization, lemmatization, stop words removal, stemming, part of speech (POS). Tokenization is the process of splitting paragraphs into sentences, or sentences into words. Capitalization brings everything to lower case for simplicity. Stop words removal removes all connecting words such as “and”, “the” or others by comparing the text to a list of stopwords. POS takes text and assigns each part of speech to each word that helps to build more understanding of a text. Stemming is a process where words are reduced to a root by removing the unnecessary suffix e.g. eating after stemming is eat. Lemmatization is an alternative approach of stemming which able to capture canonical forms based on a word’s lemma. It uses part of speech and WordNet’s lexical database of English for removing inflection. The word better, stemming will fail to provide any lemma but lemmatization would result good.

Another way to extract features from the text is Bag of Words (BoW). It is a model used in natural language processing and categorizes documents based on a vocabulary of words and occurrence of words. The common used NLP method is Vector Space Modeling (VSM). It is a way to represent documents into multidimensional space for information retrieval and documents classification and clustering. For example, if you give a query then VSM will find relevant items or term from a corpus.

II. CONTRIBUTIONS

A. Requirements Elicitation and Discovery

The manual process of requirement elicitation is expensive in terms of effort and resources. A project’s success majorly depends on the precise identification of stakeholder’s expectations and requirements for the system they desired. A possibility to do requirements elicitation is to mine available datasets e.g. social media, requirement documents, and app stores reviews etc. This mining process was performed with help of different techniques e.g. NLP and text mining technique

[21] [14]. The latest trend for identifying user requirements is to mine data obtained from platforms like Twitter, Google Play Store, and Apple Store etc. by applying ML. These user reviews are not structured requirements and contain useful information with extra information and noise that make manual requirement elicitation a difficult and challenging task. Automated requirement elicitation is helpful in these cases and can significantly reduce time, effort, and cost. This is mainly an ML *classification* task: give the set of information and identify it as a requirement or not. Sometimes *clustering* is also used for auxiliary tasks.

Guzman *et al.* [20] proposed the ALERTme approach for classifying, grouping and ranking tweets during the software evolution development. Many users shared their opinions about various software on Twitter. The huge amount of dataset made it hard to manually identify tweets that contained user requirements. The proposed methodology classified tweets as improvement requests or not, using *Naive Bayes* algorithm. This was the first study of its kind that was performed on software related tweets. The classifier was trained with following steps: 1) conversion of pre-processed tweets into a VSM model, 2) train a classifier on a set of manually annotated tweets, 3) predict the tweets categories using trained classifier. Furthermore, improvement requests were considered for the grouping which helped to sort the requests and summarize them accordingly. The summarization process contained highly ranked tweets based on parameters including likes, sentiments, and number of shares etc.

Williams *et al.* [51] performed a similar study on tweets in order to classify them as user requirements. It used basic pre-processing techniques and applied VSM on data. For the learning process, manually annotated (labelled) tweets were used and *Naive Bayes* algorithm was applied for classification. The authors claimed with the help of results that software tweets are neutral in nature, meaning sentiment analysis did not influence the outcome of the ML algorithm. It showed improved results in comparison to [20]. The study used 4000 randomly selected tweets from ten different softwares including Microsoft Visual Studio, Google Chrome, and Instagram etc.

Jiang *et al.* [24] mined user reviews from app stores for discovering evolutionary requirements. It first extracted opinions about software features from reviews. For automated opinion identification, syntactic relation based propagation approach was used that extracted targets and sentiment words iteratively using known and extracted words. Afterwards, it applied k-mean clustering for opinion categorization. The proposed system also helped developers to decide requirements related to software revenue by considering economic factors. It used two datasets of online reviews: one from the Karplersky internet security 2011 software package (from Amazon) with 380 reviews; the other one comprising 461 reviews for the TuneIn Radio Pro V3.6 mobile app (from the app store).

Lange *et al.* [29] mapped the software requirement elicitation process onto an existing military tool *skiweb*. Skiweb was used to make decisions about what actions need to be

taken on a military command. Different users posted and updated events and information using this tool. The goal of adding learning capability was to find additional information relevant to user posts. The proposed recommender system used supervised *Naive Bayes* algorithm to classify text documents in order to find related requirements to the post. Furthermore, it used topic modeling to identify the key stakeholders and suggested them requirements for further analysis according to their interest. This study used an internal organizational dataset Skiweb Data such as wiki and blogs etc.

Jha *et al.* [23] discovered user requirements by mining app store reviews. The requests were classified into three categories; *bugs*, *features*, and *junk*. The proposed methodology applied Naive Bayes and SVM. The distinction between types of sentences was identified by frame semantics *explain frame semantics* instead of text classification methods. It generated frames for each review, rather than each word. Due to small number of features, a slower dimensional model was produced with enhanced prediction capabilities. It combined existing datasets from past studies and reviews for iOS apps including CreditKarma, Fitbit, and Gmail.

Maalej presented in [34] a study on how to classify app reviews as bug reports, feature requests, user experiences, and ratings. It used *Naive Bayes* algorithm due to better results in comparison to other algorithms for classification. It also highlighted that binary classifier performed better than multi classifiers. It used meta model to enhance the classification performance e.g. ratings, tense, and sentiment scores etc. A dataset of 4400 manually annotated reviews from Google Play Store and the Apple App Store were used for the study.

Herrera *et al.* [5] built a recommender system to manage a large number of stakeholders participation in the requirements elicitation and prioritization process. In this system, stakeholders could work collaboratively to transform their needs into sets of articulated and prioritized requirements. It automatically generated specialized topics for building forums for stakeholders collaboration and discussion. Stakeholders interests were depicted from their user profiles, that also helped to create recommendations according to the interest of a community of similar stakeholders. For identifying topics, an unsupervised agglomerative clustering algorithm was applied to unstructured data. The proposed system analyzed online datasets that were gathered from stakeholders in natural language. The evaluation dataset was a collection of 36 feature requests created by graduate-level students for an Amazon-like student web-portal system.

B. Requirements Specification and Analysis

Software requirements specifications are usually stated in informal, imprecise and ambiguous natural language, thus analyzing them is a challenging task. The success of a system does not solely depend on its functional requirements. Like for functional requirements, it also significantly depends on the adherence to non-functional requirements. The primary focus of requirements analysis is generally towards the identification and specification of FRs. NFRs are usually identified and

specified in later development stages which can increase the risks of problems during the development lifecycle. Non-functional requirements may not be explicitly mentioned in a formal specification requirements documents even though they exist for all systems in freeform documents like interview notes, meeting minutes. All types of requirements are analyzed differently, and as such, it is useful to separate them. This distinction helps in managing changes in requirements as well as in precisely incorporating them in the development of the system. Manual requirement division into FRs, NFRs, and FRs and NFRs is difficult and time-consuming. Machine learning can be useful to support analysts in the error-prone task of manually discovering and classifying them which will help to ease further analysis in further processing to analyse.

1) *Identifying Non-Functional Requirements*: This is a classification problem as from a set of requirements we want to decide whether or not requirements belong to NFRs .

In a study by Slankas *et al.* [44] the authors automatically identify and classify sentences in natural language from use agreements, install manuals, regulations, requests for proposals *proposals of what?*, requirements specifications, and user manuals into 14 different NFRs categories such as *Access Control*, *Audit*, *Availability*, and *Legal* etc. Their two-step process: 1) parse natural language and turn sentences into graphs, 2) classify sentences into categories with k-nearest neighbor algorithm led them into finding 20 keywords for each category of NFRs as features for their classifier. They trained the NFR classifier with a wide variety of open and closed source EHRs (Electronic Health Record), various industry standards (HL7, CCHIT), governmental regulations, and other document sources exist to elicit documentation.

Cleland-Huang *et al.* [7] explored a similar approach and used the k-nearest neighbor classifier for grouping non-functional requirements e.g. *availability*, *look-and-feel*, *legal*. For training their classifier, they used 15 requirements specifications developed as term projects by master students at DePaul University.

2) *Identifying Functional Requirements*: This is a classification problem as from a set of requirements we want to decide whether or not requirements belong to FRs.

Wang *et al.* [49] applied a combination of machine learning, natural language processing, and semantic analysis to automatically extract functional requirements and classify them into ten different cases e.g. action, objective, goal, temporal, and constraints etc. Their framework employed techniques of semantic role labeling (which assigns a role label for each word in the sentence) and machine learning and has four steps: corpus construction, NLP preprocessing, feature extraction and EFRF (Extended Functional Requirements Frame) functional cases extraction. The authors trained their bi-directional LSTM-CRF network, a variant of Recurrent Neural Networks architecture model, with an E-commerce requirements dataset and tested it on requirements of automaker systems. They proposed EFRF through analyzing the linguistic characterization of software requirement specifications. EFRF consists of the 10 FRs

types mentioned above and allows capturing the semantic information in the natural language.

3) *Distinguishing Functional from Non functional Requirements*: This is a classification problem as from a set of requirements we want to decide whether or not requirements belong to a certain class.

Lu *et al.* [30] automatically classified text from user reviews for (app) stores into FR, NFR, and others. The authors further classified NFRs into four categories including reliability, usability, portability, and performance. The approach used a supervised machine learning algorithm *bagging*. The text trimming used standard pre-processing steps. Also, the sentences were augmented by several most similar words to the user reviews in the training set and showed that augmented user reviews can lead to better classification results. The machine learning algorithm. Bagging is more suitable for NFRs classification from user reviews than Naïve Bayes and J48. The authors used 6696 raw user reviews from iBooks and 4400 raw user reviews from WhatsApp as datasets.

Deoxadez *et al.* [11] used semi-supervised classification techniques for automated classification of FRs and NFRs from user reviews from the app store. This study dealt with two problems: 1) minimizing the annotation effort for labeling the big dataset of user reviews, and 2) classification of FRs and NFRs. The proposed solution to the first problem used the semi-supervised self-labeling algorithm. Self-labeling algorithms required a small amount of dataset to get comparable results as supervised techniques. Features were obtained by applying the standard pre-processing and BoW algorithm. For user reviews classification *Naive Bayes* algorithm was selected, because of its general high-performance results in classification problems. It used reviews of top 40 paid and free apps from ten different categories.

Kutranvoic *et al.* [28] performed automated analysis on software requirement documents written in natural language and classify them into FRs, NFRs and subcategories of NFRs. The basic pre-processing steps were applied on requirements before applying the ML algorithm. It used a supervised support vector machine (SVM) algorithm for classification. Some NFRs were ignored because of their minor presence in the documents. Also, the dataset was imbalanced that meant it was not equally distributed. For avoiding data imbalance problem additional dataset i.e. user comments from Amazon was integrated into the main dataset. The basic pre-processing steps were applied on requirements before applying the ML algorithm. The study used data from open source tera PROMISE repository that consists of 625 labeled natural language requirements (255 FRs [40.8%] and 370 NFRs [59.2%]).

Abad *et al.* [1] targeted two similar problems: the first one is the classification of FR and NFR, and the second classification of NFR into categories. It first performed pre-processing for text trimming. Afterward, increased the weight of influential words in the dataset using feature co-occurrence and regular expression. It learnt J48 decision tree (DT) for classification of FRs and NFRs. It used Binarized Naive Bayes

(BNB) for sub-classification of NFRs. BNB outperformed among other algorithms such as clustering, k-means, and hybrid clustering and k-means. The study showed that the text pre-processing approaches positively influenced on classification. It improved FRs and NFRs classification accuracy from 89.92% to 95.04%. It used reviews of top 40 paid and free apps from ten different categories.

Garzoli [17] proposed a system for requirement analysis that allowed to identify software functionalities within large collections of requirements written in natural language for requirement analysis. It classified the large dataset into five types: FRs, NFRs, *design and construction constraints*, *operator requirements* and *performance requirements*. The goal of the study was to come up with a general architecture for large-scale and adaptive requirement analysis. It used BoW with text mining techniques for lexical and grammatical feature for information retrieval and SVM model for classification of requirements. The learning classifiers contained 4,727 annotated requirements, related to three different scenarios. The dataset was taken from the combat management system of a Naval Combat System. The learning classifiers contained 4,727 annotated requirements, related to three different scenarios

Wieloch *et al.* [50] introduced methods that enhance the accuracy of tracing (classifying) requirements which occurred frequently across multiple projects and/or domains. In the paper, they presented Trace by Classification (TBC), a machine learning approach in which a classifier is trained to identify and classify requirements and/or other kinds of software artifacts which occur relatively frequently across different projects (they call this process generating trace links for software artifacts in their research). Then the training phase identified a set of indicator terms for each NFR category and the classifier trained by the set of identified weighted indicator terms that can be used for the last step which is to classify additional artifacts into functional and non-functional (e.g. look-and-feel, performance, security, etc). A probability value represented a possibility that the new requirement belonged to a certain artifact type computed as a function of the occurrence of indicator terms of that type in the requirement.

4) *Requirement Prioritization*: Complex software system generally has thousands of requirements with multiple stakeholders and customers. Each one of them has their own set of requirements and opinions and wants their requirements implementation accordingly. However, several factors make implementation process of all the requirements infeasible and hard e.g. budget, and different opinion among stakeholders etc. Therefore, it is important to make a proper decision for prioritizing requirements considering all the factors for the success of the project. Different models exist in the literature for prioritization of software requirements such as analytical hierarchical process (AHP) [43], Goal oriented [48], cost value approach [26] etc. In these techniques, human input is majorly involved. Qaddoura *et al.* [38] reviewed the prioritization techniques and also shed light on the ML contribution. ML can be used for automated analysis of these large set of software requirements prioritization, and it can also help to improve the

existing techniques.

Dhingra *et al.* [12] predicted the most appropriate technique for software requirement prioritization process. The input from the user was taken as characteristic values (detect consistency, maintain information, not available, or both) for different attributes. These attributes list included consistency, traceability, priority basis, rigorous/systematic, distributed stakeholder, cognitive aspects, and human experience. The output was the most appropriate requirement prioritization method e.g. AGORA, AHP etc. Fuzzy interface technique was used for prediction of prioritization method based on certain rules. The proposed system framework had three phases; training phase, fuzzing inference process, and testing phase. The drawback of fuzzy approach was the wrong prediction for boundary values, which was overcome by adopting decision trees. DT learned from datasets and predicted the most suitable prioritization technique. Out of 45 test samples, the framework was able to classify 43 tests accurately.

Avesani *et al.* [3] presented the study that dealt with the scalability problems which arise in managing the prioritization of a large number of requirements specified in the the AHP technique. The existing solution to scalability issues used heuristics. It helped to decide when to stop the pairwise elicitation process. The proposed framework outperformed AHP by giving an accurate approximation of the final ranking within a limited elicitation effort. It used rank-based learning algorithm and produced a ranking of all requirements. This technique built up the solution by looking at examples. The input for the learning algorithm were a finite set of requirements, the ranking criteria, initial user preferences and density function.

A similar study was performed in [4] from the same group [3] for identification of decision-making issues related to the management of risks in Open Source Software adoption in medium and large organizations. A semi-automated system was proposed that used case based ranking classification algorithm. The input was priority elicitation of goals by the decision maker and risk goal ranking function (predefined ranking criteria ordering the goal). As an output, it ranked the final risk-based goals.

5) *Security Requirements*: Due to the orthogonal character of their impact on a system, *security* requirements are notoriously difficult to identify, objectify and quantify [39]. Also during requirement specification, it very often happens that security requirements are masked by functional requirements (but can be deduced from the context of the domain the system operates in) [42]. Because of this, it often happens in practice that security requirements are often only marginally tackled during system construction, paving the way to potentially catastrophic consequences. Machine learning can be of use here by aiding in the identification of segments of text that describe security requirements. This is a *classification* problem: given a text, identify which parts of it correspond to which type of security issues.

Jindalet *et al.* [25] automatically learn decision trees that can be used to classify security requirements as *authentication*, *access control*, *encryption* or *data integrity*. Preprocessing of

the data is done by stemming relevant terms and the *features* used are such terms.

Riaz and her colleagues [42] use the k-nearest neighbors algorithm to classify sentences in requirements documents as *confidentiality*, *integrity*, *authentication*, *availability*, *accountability* or *privacy* requirements. In order to find adequate sentences and provide context to the classifier, the authors start by finding a type for each sentence among the possibilities *title*, *list start*, *list element* or *normal sentence*. For the classification the authors use the number of word transformations needed to go from one term in one sentence to a term in another sentence. The classifier is trained using requirements sentences from the healthcare domain that are manually classified. A particularity of the approach is that each security requirement type is associated to a template that helps in translating the security requirements into functional requirements in order to ease during the implementation of the final system.

C. Requirements Validation

1) *Traceability*: Validation is to guarantee that requirements are reflecting stakeholders' needs, confirm the quality of the system, consistency, and traceability. In requirements traceability, the emphasis is on the ability to track the life of requirements and their established links within other artifacts. However, the main barrier assures traceability is the needed effort for building and maintaining the links between those artifacts. That is why many research has tried to apply machine learning and automated tools for facilitating the establishment of links [18]. Traceability tackled in the research mainly by the use of machine learning classification and reinforcement learning methods.

Gervasi *et al.* [18] investigate what can be learned from links that are already established. They build classifiers as a mean to develop models of tracing that can then be interpreted by humans to understand how requirement tracing is done in practice. Their purpose is to revise the existing models of hard-coded traceability tools such as VSM. They used a publicly-available dataset of requirements with traceability information, originally based on the CM-1 project by the NASA Metrics Data Program. Their approach has the following steps: applied preprocessing techniques and transferred requirements into a vector of features which from them derived set of classification cases by joining one high-level requirement and one low-level requirement and adding a classification of *link* or *nolink* based on whether that particular pair was a true link in the original dataset, or not. Finally, use the dataset to train and test two different classifiers from the WEKA collection, a Naive Bayesian classifier, and the J48 decision-tree classifier.

Sultanov *et al.* [46] finds traceability candidates from high-level to low-level requirements by the use of reinforcement learning. They used textual high and low-level requirements documents as an input and try to find the candidate traces. Their technique demonstrated statistically significantly better results than the Information Retrieval technique.

D. Requirements Management

1) *Visualization*: Natural language requirement documents can be hard to comprehend and analyze. Similarly, stakeholders have to review and understand requirements for large and complex systems. In these scenarios, basic information visualizations, like charts and graphs have been used in requirements engineering. These visualizations are usually applied to improve textual requirements with summarization that combined large amounts of information into a single representation for quick consumption by stakeholders [41]. Machine learning is useful in discovering visualized groups of large numbers of requirements. In research both clustering and classification methods were used for this purpose.

ReCVisu (Requirements Clustering Visualization) tool is presented in Reddivari *et al.* [41] paper. ReCVisu, an exploration tool based on quantitative visualizations helps requirements engineers understand the nature of the requirements in a visual form. In ReCVisu, the dependence graph consists of requirements artifacts as nodes and the textual similarities as edges. The automatic grouping of requirements into clusters can help in areas such as uncovering the requirements structure, navigating around the requirements space, modularizing crosscutting concerns, and understanding requirements interactions and evolution.

Pinqui *et al.* [37] recognize the enormous volume of requirements as big data with which companies struggle to make strategic decisions early on. Therefore, they built a complete visual framework to filter requirements from stakeholders in a way that architects can make better insightful decisions. They suggest training a multi-class SVM model from domain-specific (mechanics, electronics, etc.) dictionaries and handbooks. Overall, the paper proposes a framework to go from management-oriented requirements to architecture-oriented requirements in which SVM is only applied in a small part of it.

Lucassen *et al.* [31] introduced an automated method for visualizing requirements at different levels of granularity. Their visualization method from user stories consists of 1) the generation of an overview which provides a general context for understanding the dataset. It used Word2Vec and Ward's clustering algorithm to build up inter cluster relationship matrix of concepts from the dataset. 2) zooming in and out mechanisms and 3) filtering techniques to reduce the complexity of the data presentation. Possible anticipated applications of this visualization are: discovering missing relationships between clusters that may result in further user stories, teaching system functionality by exploring simplified, manageable chunks, and analyzing expected system changes after introducing new sets of user stories.

2) *Structuring Documents*: Requirements of the system are usually presented in natural language documents. These documents require to be properly structured for a better overall understanding of the requirements. For this purpose, the document should be organized with independent sections which each one contains conceptually connected requirements [16].

Moreover, technical review is a usual way to guarantee the quality in natural language specifications. However, extensive and comprehensive specifications make it problematic for reviewers to find defects, especially consistency or completeness ones. Therefore, use of machine learning algorithms can support reviewers with their work by automatically classifying and clustering the information that is spread over many sections of many documents [36].

Duan *et al.* [15] used hierarchical automated clustering technique for detecting cross-cutting concerns as it is beneficial for the process of requirements analysis and architectural design. The reported experiments in this paper were supported by two tool sets, Poirot, a web-based tool designed to generate traces between various software engineering artifacts which was applied to compute similarity scores between requirements and a developed prototype tool, capable of reading structured requirements specification and generated similarity scores and then clustering requirements.

Winkler *et al.* [52] applied convolutional neural networks to automatically classify content elements of a natural language requirements specification as "requirement" or "information". Their approach increases the quality of requirements specifications as it distinguishes important content for activities. For converting natural language into a vector representation word2vec method is used. A set of 10000 content elements extracted from 89 requirements specifications of an industry partner used for training the network through the use of Tensorflow library using stochastic gradient descent.

Ferrari *et al.* [16] automatically recognize the sections in the document that need requirements relatedness and sections independence that help to enhance the document structure. The authors defined a novel algorithm named Sliding Head-Tail Component (S-HTC) for clustering the requirements according to relatedness (the algorithm is based on known distance - Jaccard similarity metric, Levenshtein distance, and the convex combination between σ_{jac} and σ_{lev}). The algorithm groups together similar requirements that are contiguous in the requirements document. The effectiveness of the algorithm was evaluated with a test on requirements standard of a railway domain (583 requirements).

Based on Rauf *et al.* [40] software specification documents usually contain instances of logical structures, such as business rules, use cases, and functional requirements. Automated identification and extraction of these instances will benefit requirements management features, like automated traceability, template conformance checking, and guided editing. The authors planned a framework that gets requirements documents as an input and tries to develop a template for the general structure of it by specifying logical structures in terms of their content, textual rendering, and variability and then the extracting the instances of such structures from rich-text documents.

Ott *et al.* [36] automatically classified and extracted requirements with related information which are spread over many sections over many documents by the use of Multinomial Naive Bayes and Support Vector Machines classification algo-

gorithms as it will be helpful for reviewers with their work. As their input, they have used two German automotive specifications (Mercedes-Benz) which describe the functional and non-functional requirements of a Doors Closure Module (DCU). A specification and its referenced documents often sum up to 3,000 pages at Mercedes-Benz. Their method collects requirements of related information into classes, which they call topic landscape and later they built a tool, ReCaRe(Review with Categorized Requirements) which is the realization of the topic landscape based on eclipse with a data connection to IBM Rational DOORS.

III. DISCUSSION

Our survey work implicitly points to a number of trends that we will concretize and summarize in this section. Note that while the pointers we provide here are informed by the literature review we conducted, this survey is not fully systematic (as described in section IV) which means our conclusions may be revised and/or extended by future surveys of the domain.

Table `tabletab:summary` summarizes our findings. It provides partial answers to **RQ2** (“What types of learning methods are used when ML is applied to RE?”) and **RQ3** (“Which are the RE problems that are currently using ML methods?”) in columns *ML Task* and *themes*, respectively. The table also provides partial responses to **RQ1** (“What is the current state of the practice in ML & RE?”) and **RQ4** (“Is using ML methods improving RE?”). The answer to **RQ1** seems to be “at its beginning”, given the prevalent lack of comparison with the state of the art as can be observed in `tabletab:summary`. The answer to **RQ4** is “unknown”, given that most of the studies read by us were initial proposals with little academic or industrial validation in real software engineering tools or projects.

Note that `tabletab:summary` provides additional information on which types of algorithms are used for each kind of theme, as well as datasets used for learning and which are available online.

It is obvious from our survey that NLP techniques are heavily used throughout a majority of the research tackling the application of ML to RE. This is not surprising and even intuitive. RE is the area of software engineering where natural language is employed more ubiquitously, as RE techniques and tools play the role of interface between stakeholders such as clients, certification entities, architects or developers. Although many attempts have been done to bring formality to requirements engineering [32], [47], the *de facto* language between technical and non-technical stakeholders for real-world projects continues being natural language, and in particular English. The IBM Rational DOORS family [22] of tools is an example of a natural-language based tool for requirement engineering that has become the reference in many domains. In the techniques we have observed, NLP is heavily used for the preprocessing stages of natural language in order to bring the data to a format that can be consumed by a learning algorithm (see section I-C).

The authors of the articles we have processed in our survey of point to the idea that ML can potentially bring about enormous benefits in terms of processing and taking decisions based on large amounts of imprecise and ambiguous data. In the real-world of software engineering, parsing and summarizing requirements is a very time-consuming activity. Also, decisions taken by technical stakeholders are often based on imprecise, incomplete and noisy information and are supported by rules-of-thumb, experience and intuition. ML is by nature built to handle and cope with such challenges – it based on data and it’s main purpose is exactly to build model of patterns that humans associate to rules-of-thumb, experience or intuitions. Additionally, ML methods often provide a precise degree of certainty regarding the correctness of decisions taking during a software engineering project. Such measures, although valid regarding the quality of the learning process, allow assessing the risk associated to certain steps in the course of a project.

In the sequence of the previous paragraph, a large set of datasets on RE are available online. We have identified a few of such datasets in `tablesummary`. This fact is a cornerstone for the domain, as most ML algorithms existing nowadays are very data-intensive. One of the authors of this survey has recently written a similar article on the application of ML to formal verification [2], for which the datasets available to learn from are typically very small and almost never made public. The authors of the article recognize that such scarceness of data is partly due to the niche nature of the domain of formal verification, where the datasets are mostly in the form of mathematical proofs. Nonetheless, and in spite of the large body of work regarding the application of ML to formal verification, such scarceness of data poses a problem not only to the automated learning, but also to the scientific validation of such proposals. This is not the case in RE, where many datasets are publicly available on which both learning and validation can be done.

The majority of the articles we found on the topic of ML and RE have to do with either the *elicitation* or the *analysis* phases of RE. These findings are compatible with the idea that parsing requirements texts and classifying the information that is contained in them is strenuous for humans and thus it is desirable that such tasks are as automatic as possible. The *validation* and *management* phases in RE also imply tasks that can be automated as we have shown through our survey, but the state of the art in the domain seems to imply that the first two phases have priority for researchers and practitioners.

Also, we have observed through our readings that while *classification* is the most used ML task, *clustering* also plays an important role in the domain of ML applied to RE. This contradicts the results in [2], where *clustering* has almost no expression in work that applied ML to formal verification. We believe this provides support to the thesis that ML is particularly appropriate to RE– given *clustering* is especially useful when mining non-formal data such as free-form text.

	Themes	Contributions	ML Task	ML Model Types	Datasets Used
E	External	[20](o) [51](+) [24](+) [29](o) [23](+) [5](o)	Classification Clustering	(Multinomial) Naïve Bayes; Support Vector Machines	Online reviews for KIS 2011 (from Amazon) and TuneIn 3.6 (from App Store) Skiweb data
	Non-Functional	[44](o) [7](o)	Classification	k-Nearest Neighbors	Open Source PROMISE Dataset ²
	Functional	[49](o)	Classification	Bi-Directional Long Short-Term Memory Conditional Random Field Network	-
	Functional & Non-Functional	[30](o) [11](o) [28](o) [1](o) [17](o) [50](o)	Classification	Bagging, Naïve Bayes, SVM	Open Source PROMISE Dataset app-store reviews
S	Prioritization	[12](o) [3](+) [4](o)	Classification	Case Based Ranking J48 DT -	-
	Security	[25](o) [42](o)	Classification	Decision-Tree k-Nearest Neighbors	-
V	Traceability	[18](o) [46](o)	Classification Reinforcement Learning	Naïve Bayes / J48 Decision-Tree	Open Source CM-1 NASA project ³ Open Source Pine Dataset ⁴
M	Visualization	[41](o) [37](o) [31](o)	Classification Clustering	Support Vector Machines Ward's method	-
	Structuring	[15](o) [52](o) [40](o) [16](o) [36](o)	Classification Clustering	Multinomial Naïve Bayes Support Vector Machines Convolutional Neural Networks Sliding Head-Tail Componen Clustering Hierarchical Clustering	International Union of Railways (EIRENE Functional Requirements Specification ⁵) Mercedes-Benz car development Specifications ⁶

Legend: (+) improves the state of the art; (-) comparable to or worse than state of the art; (o) no information on how the approach relates to the state of the art

TABLE I: Contributions and ML tasks related to each theme within each RE approach.

IV. THREATS TO VALIDITY

The validity of the study might be affected by the coverage of the search results, bias on study selection, and inaccuracy of data extraction.

Study Coverage: The study we present here is partial, meaning relevant work could be missing due to inadequate search strings or the list of databases not being complete. The data preparation was based as much as possible on a systematic method, which resulted on a map of the read articles and their main features as relevant to our study.

Study Selection Bias: we understand that the assessment might be biased by the interests of the involved researchers. As such, the themes that we discuss in this article may be influenced by the preferences of the involved researchers. To mitigate for this threat, a set of include and exclude criteria was predefined and researchers assessed the title and abstract of the papers to steer the research. For many of the papers retrieved by our queries apply NLP to requirements engineering but involve no learning (e.g. [53], [10], [6], to cite a few). We have explicitly excluded such papers from our survey: although NLP tools do sometimes include ML algorithms, their functionality is used in a black-box manner by RE researchers and as such were not taken into consideration by our work.

Inaccuracy of Data Extraction: the data extraction process might be biased by researcher interest. To mitigate for this threat, the selection of data items was strictly driven by the research questions. Moreover, reading assignments were marked by the researchers depending on their confidence level. Low-confidence assignments were discussed between the authors until a consensus was reached.

V. CONCLUSION

Through our bird's eye view of ML applied to RE we have observed that in the past couple of decades a good amount of research has been done on how to bring these two worlds together. The stakes are high: while requirements engineering is currently a domain under intensive research, attempts to address its challenges academically have translated into few results in practice. Free-form text-based tools with light-weight structuring capabilities such as DOORS are now the norm in practice. Requirements elicitation, analysis, validation and management keep on relying on human expertise and talent. While academics often insist that better formalization brings advantages, the languages in which requirements are formalized do not match the need that stakeholders in the RE process (technical and non-technical) need to communicate through artifacts that are intelligible to all.

While not overstating the potential of ML, which has its own challenges to overcome such as coarseness of the learned models, overfitting or hungriness for data, we have provided in this article indications that ML might become a cornerstone in RE. For now, it seems like the domain is now undergoing a pre-scientific phase: the studies we have analyzed seldom compare themselves with the state-of-the-art (see tablesummary). This suggests that the current body of research is composed of new ideas, which have not yet been validated to its full extent by the scientific or industrial communities. We thus call for a more extensive survey to validate the preliminary conclusions we present in this work.

REFERENCES

- [1] Z. S. H. Abad, O. Karras, P. Ghazi, M. Glinz, G. Ruhe, and K. Schneider. What works better? A study of classifying requirements. *CoRR*, abs/1707.02358, 2017.
- [2] M. Amrani, L. Lúcio, and A. Bibal. ML + FV = Love? A survey on the application of machine learning to formal verification. *CoRR*, abs/1806.03600, 2018.
- [3] P. Avesani, C. Bazzanella, A. Perini, and A. Susi. Facing scalability issues in requirements prioritization with machine learning techniques. In *13th IEEE International Conference on Requirements Engineering (RE'05)*, pages 297–305, Aug 2005.
- [4] P. Avesani, A. Perini, A. Siena, and A. Susi. Goals at risk? machine learning at support of early assessment. In *2015 IEEE 23rd International Requirements Engineering Conference (RE)*, pages 252–255, Aug 2015.
- [5] C. Castro-Herrera, C. Duan, J. Cleland-Huang, and B. Mobasher. A recommender system for requirements elicitation in large-scale software projects. In *Proceedings of the 2009 ACM Symposium on Applied Computing, SAC '09*, pages 1419–1426, New York, NY, USA, 2009. ACM.
- [6] H. Chen, K. He, P. Liang, and R. Li. Text-based requirements preprocessing using nature language processing techniques. In *2010 International Conference On Computer Design and Applications*, volume 1, pages V1–14–V1–18, June 2010.
- [7] J. Cleland-Huang, R. Settini, X. Zou, and P. Solc. Automated classification of non-functional requirements. *Requirements Engineering*, 12(2):103–120, Apr 2007.
- [8] J. Coughlan and R. D. Macredie. Effective communication in requirements elicitation: A comparison of methodologies. *Requir. Eng.*, 7(2):47–60, June 2002.
- [9] A. M. Davis. *Software Requirements: Objects, Functions, and States*. Prentice-Hall, Inc., Upper Saddle River, NJ, USA, 1993.
- [10] D. K. Deepthimahanti and R. Sanyal. Semi-automatic generation of uml models from natural language requirements. In *Proceedings of the 4th India Software Engineering Conference, ISEC '11*, pages 165–174, New York, NY, USA, 2011. ACM.
- [11] R. Deocadez, R. Harrison, and D. Rodriguez. Automatically classifying requirements from app stores: A preliminary study. In *2017 IEEE 25th International Requirements Engineering Conference Workshops (REW)*, pages 367–371, Sept 2017.
- [12] S. Dhingra, S. G. M. Madan, and M. R. Selection of prioritization technique for software requirement using fuzzy logic and decision tree. In *2016 Online International Conference on Green Engineering and Technologies (IC-GET)*, pages 1–11, Nov 2016.
- [13] P. Domingos. *The Master Algorithm: How the Quest for the Ultimate Learning Machine Will Remake Our World*. Basic Books, 2015.
- [14] L. Dong, X. Zhang, N. Ye, and X. Wan. Research on user requirements elicitation using text association rule. In *Intelligence Information Processing and Trusted Computing (IPTC), 2010 International Symposium on*, pages 357–359. IEEE, 2010.
- [15] C. Duan and J. Cleland-Huang. A clustering technique for early detection of dominant and recessive cross-cutting concerns. In *Proceedings of the Early Aspects at ICSE: Workshops in Aspect-Oriented Requirements Engineering and Architecture Design, EARLYASPECTS '07*, pages 1–, Washington, DC, USA, 2007. IEEE Computer Society.
- [16] A. Ferrari, S. Gnesi, and G. Tolomei. Using clustering to improve the structure of natural language requirements documents. In J. Doerr and A. L. Opdahl, editors, *Requirements Engineering: Foundation for Software Quality*, pages 34–49, Berlin, Heidelberg, 2013. Springer Berlin Heidelberg.
- [17] F. Garzoli, D. Croce, M. Nardini, F. Ciambra, and R. Basili. Robust requirements analysis in complex systems through machine learning. In A. Moschitti and B. Plank, editors, *Trustworthy External Systems via Evolving Software, Data and Knowledge*, pages 44–58, Berlin, Heidelberg, 2013. Springer Berlin Heidelberg.
- [18] V. Gervasi and D. Zowghi. Mining requirements links. In D. Berry and X. Franch, editors, *Requirements Engineering: Foundation for Software Quality*, pages 196–201, Berlin, Heidelberg, 2011. Springer Berlin Heidelberg.
- [19] M. Glinz. On non-functional requirements. In *15th IEEE International Requirements Engineering Conference (RE 2007)*, pages 21–26, Oct 2007.
- [20] E. Guzman, M. Ibrahim, and M. Glinz. A little bird told me: Mining tweets for requirements and software evolution. In *2017 IEEE 25th International Requirements Engineering Conference (RE)*, pages 11–20, Sept 2017.
- [21] C. Hollis and T. Bhowmik. Automated support to capture verbal just-in-time requirements in agile development: A practitioner view. In *2017 IEEE 25th International Requirements Engineering Conference Workshops (REW)*, pages 419–422, Sept 2017.
- [22] IBM. The IBM Rational DOORS family. <https://www.ibm.com/us-en/marketplace/rational-doors>.
- [23] N. Jha and A. Mahmoud. Mining user requirements from application store reviews using frame semantics. In P. Grünbacher and A. Perini, editors, *Requirements Engineering: Foundation for Software Quality*, pages 273–287, Cham, 2017. Springer International Publishing.
- [24] W. Jiang, H. Ruan, L. Zhang, P. Lew, and J. Jiang. For user-driven software evolution: Requirements elicitation derived from mining online reviews. In V. S. Tseng, T. B. Ho, Z.-H. Zhou, A. L. P. Chen, and H.-Y. Kao, editors, *Advances in Knowledge Discovery and Data Mining*, pages 584–595, Cham, 2014. Springer International Publishing.
- [25] R. Jindal, R. Malhotra, and A. Jain. Automated classification of security requirements. In *2016 International Conference on Advances in Computing, Communications and Informatics (ICACCI)*, pages 2027–2033, Sept 2016.
- [26] J. Karlsson and K. Ryan. A cost-value approach for prioritizing requirements. *IEEE Softw.*, 14(5):67–74, Sept. 1997.
- [27] G. Kotonya and I. Sommerville. *Requirements Engineering: Processes and Techniques*. Wiley Publishing, 1st edition, 1998.
- [28] Z. Kurtanović and W. Maalej. Automatically classifying functional and non-functional requirements using supervised machine learning. In *2017 IEEE 25th International Requirements Engineering Conference (RE)*, pages 490–495, Sept 2017.
- [29] D. S. Lange. Text classification and machine learning support for requirements analysis using blogs. In *Innovations for Requirement Analysis. From Stakeholders' Needs to Formal Designs*, pages 182–195, Berlin, Heidelberg, 2008. Springer Berlin Heidelberg.
- [30] M. Lu and P. Liang. Automatic classification of non-functional requirements from augmented app user reviews. In *Proceedings of the 21st International Conference on Evaluation and Assessment in Software Engineering, EASE'17*, pages 344–353, New York, NY, USA, 2017. ACM.
- [31] G. Lucassen, F. Dalpiaz, J. M. E. M. van der Werf, and S. Brinkkemper. Visualizing user story requirements at multiple granularity levels via semantic relatedness. In I. Comyn-Wattiau, K. Tanaka, I.-Y. Song, S. Yamamoto, and M. Saeki, editors, *Conceptual Modeling*, pages 463–478, Cham, 2016. Springer International Publishing.
- [32] L. Lúcio, S. Rahman, C.-H. Cheng, and A. Mavin. Just formal enough? automated analysis of ears requirements. In *NASA Formal Methods - 9th International Symposium, NFM 2017, Moffett Field, CA, USA, Proceedings*, volume 10227 of *Lecture Notes in Computer Science*, pages 427–434. Springer, 2017.
- [33] T. M. Mitchell. *Machine Learning*. McGraw-Hill, 1997.
- [34] W. Maalej and H. Nabil. Bug report, feature request, or simply praise? on automatically classifying app reviews. In *2015 IEEE 23rd International Requirements Engineering Conference (RE)*, pages 116–125, Aug 2015.
- [35] B. Nuseibeh and S. Easterbrook. Requirements engineering: A roadmap. In *Proceedings of the Conference on The Future of Software Engineering, ICSE '00*, pages 35–46, New York, NY, USA, 2000. ACM.
- [36] D. Ott. Automatic requirement categorization of large natural language specifications at mercedes-benz for review improvements. In J. Doerr and A. L. Opdahl, editors, *Requirements Engineering: Foundation for Software Quality*, pages 50–64, Berlin, Heidelberg, 2013. Springer Berlin Heidelberg.
- [37] R. Piquié, P. Véron, F. Segonds, and N. Croué. A collaborative requirement mining framework to support oems. In Y. Luo, editor, *Cooperative Design, Visualization, and Engineering*, pages 105–114, Cham, 2015. Springer International Publishing.
- [38] R. Qaddoura, A. Abu-Srhan, M. H. Qasem, and A. Hudaib. Requirements prioritization techniques review and analysis. In *2017 International Conference on New Trends in Computing Sciences (ICTCS)*, pages 258–263, Oct 2017.
- [39] A. Rashid, S. A. A. Naqvi, R. Ramdhany, M. Edwards, R. Chitchyan, and M. A. Babar. Discovering security requirements. In *2016 IEEE/ACM 38th International Conference on Software Engineering (ICSE)*, pages 866–876, May 2016.

- [40] R. Rauf, M. Antkiewicz, and K. Czarnecki. Logical structure extraction from software requirements documents. In *2011 IEEE 19th International Requirements Engineering Conference*, pages 101–110, Aug 2011.
- [41] S. Reddivari, Z. Chen, and N. Niu. Recvisu: A tool for clustering-based visual exploration of requirements. In *2012 20th IEEE International Requirements Engineering Conference (RE)*, pages 327–328, Sept 2012.
- [42] M. Riaz, J. King, J. Slankas, and L. Williams. Hidden in plain sight: Automatically identifying security requirements from natural language artifacts. In *2014 IEEE 22nd International Requirements Engineering Conference (RE)*, pages 183–192, Aug 2014.
- [43] T. L. Saaty. Decision making with the analytic hierarchy process. *International journal of services sciences*, 1(1):83–98, 2008.
- [44] J. Slankas and L. Williams. Automated extraction of non-functional requirements in available documentation. In *2013 1st International Workshop on Natural Language Analysis in Software Engineering (NaturaLiSE)*, pages 9–16, May 2013.
- [45] C. Standish Group. The standish group report, 2014.
- [46] H. Sultanov and J. H. Hayes. Application of reinforcement learning to requirements engineering: requirements tracing. In *2013 21st IEEE International Requirements Engineering Conference (RE)*, pages 52–61, July 2013.
- [47] S. Teufl. *Seamless Model-based Requirements Engineering: Models, Guidelines, Tools*. PhD thesis, Technical University Munich, Germany, 2017.
- [48] A. Van Lamsweerde. Goal-oriented requirements engineering: A guided tour. In *Proceedings of the Fifth IEEE International Symposium on Requirements Engineering*, RE '01, pages 249–, Washington, DC, USA, 2001. IEEE Computer Society.
- [49] Y. Wang and J. Zhang. Experiment on automatic functional requirements analysis with the efrfs semantic cases. In *2016 International Conference on Progress in Informatics and Computing (PIC)*, pages 636–642, Dec 2016.
- [50] M. Wieloch, S. Amornborvornwong, and J. Cleland-Huang. Trace-by-classification: A machine learning approach to generate trace links for frequently occurring software artifacts. In *2013 7th International Workshop on Traceability in Emerging Forms of Software Engineering (TEFSE)*, pages 110–114, May 2013.
- [51] G. Williams and A. Mahmoud. Mining twitter feeds for software user requirements. In *2017 IEEE 25th International Requirements Engineering Conference (RE)*, pages 1–10, Sept 2017.
- [52] J. Winkler and A. Vogelsang. Automatic classification of requirements based on convolutional neural networks. In *2016 IEEE 24th International Requirements Engineering Conference Workshops (REW)*, pages 39–45, Sept 2016.
- [53] X. Xiao, A. Paradkar, S. Thummalapenta, and T. Xie. Automated extraction of security policies from natural-language software documents. In *Proceedings of the ACM SIGSOFT 20th International Symposium on the Foundations of Software Engineering*, FSE '12, pages 12:1–12:11, New York, NY, USA, 2012. ACM.
- [54] P. Zave. Classification of research efforts in requirements engineering. *ACM Comput. Surv.*, 29(4):315–321, Dec. 1997.
- [55] D. Zhang and J. J. P. Tsai. Machine learning and software engineering. In *14th IEEE International Conference on Tools with Artificial Intelligence, 2002. (ICTAI 2002). Proceedings.*, pages 22–29, 2002.
- [56] D. Zowghi and C. Coulin. *Requirements Elicitation: A Survey of Techniques, Approaches, and Tools*, pages 19–46. Springer Berlin Heidelberg, Berlin, Heidelberg, 2005.